

# TSA Assignment 9

Christian Hilscher - 1570550

18/5/2020

## Question 15

### DGP

For gathering the data the following function is used. The *models* input is a list containing strings to identify the models as given in the assignment. It first generates the errors drawn from different distributions and then uses the function *generate\_x* to provide the finished time series. The results are then stored in a dictionary with the keys being the individual model names.

```
DGP <- function(tslength, samplesize, b, models){

  storage <- dict()
  for (model in models){
    storage[[model]] <- vector()

    # Drawing errors
    for (i in seq(1, samplesize)){
      if (model=="i"){
        errors <- rnorm(1001)
      }else if (model=="ii"){
        errors <- rt(1001, 5)
      }else if (model=="iii"){
        errors <- runif(1001)
      }else{
        stop('Please provide one of the models: i, ii, iii')
      }

      # Getting the time series with the corresponding errors
      storage[[model]] <- rbind(storage[[model]],
                               generate_x(tslength, errors, b))
    }
  }
  return(storage)
}

generate_x <- function(tslength, errors, b){

  e_shifted <- shift(errors,1)
  x <- (errors+b*e_shifted)[2:1001]
  return(x)
}
```

For setting up the whole process I run the following specifications. Here I also provide the true variances of

the individual models because I need them later on. The vector containing the  $\lambda_j$  made according to the slides, nothing fancy there.

```
TS_length <- 1000
Sample_size <- 1000
b <- 0.25
models <- cbind('i', 'ii', 'iii')
sigma_list <- cbind(1, 5/3, 1/12)

data <- DGP(TS_length, Sample_size, b, models)
M <- floor((TS_length-1)/2)
lambdas <- (2*pi/TS_length) * seq(1:M)

# Making space for end results
df_final <- data.frame()
```

## Frequency Domain

### Theoretical Setup

From the  $MA(1)$  model  $x_t = \epsilon_t + b\epsilon_{t-1}$  I can infer the following properties:

$$I_x(\lambda_j) = \frac{1}{2\pi T} \left( \left( \sum_{t=1}^T x_t \cos[\lambda_j(t-2)] \right)^2 + \left( \sum_{t=1}^T x_t \sin[\lambda_j(t-2)] \right)^2 \right)$$

$$f_x(\lambda) = (1 + b^2 + 2b\cos(\lambda)) \frac{\sigma^2}{2\pi}$$

$$I_M(b, \sigma^2) = \sum_{j=1}^M \left[ -\ln \left( (1 + b^2 + 2b\cos(\lambda)) \frac{\sigma^2}{2\pi} \right) - \frac{I_x}{(1 + b^2 + 2b\cos(\lambda)) \frac{\sigma^2}{2\pi}} \right] \quad (*)$$

The aim is to maximise (\*) for which I first take the derivative and single out  $\sigma^2$  and plug it back into (\*) such that I only have to optimize over  $b$ .

Taking the derivative, setting it equal to zero and solving for  $\sigma^2$  yields:

$$\begin{aligned} \frac{\partial I_M}{\partial \sigma^2} &= \sum_{j=1}^M -\frac{1}{\Delta \frac{\sigma^2}{2\pi}} \Delta 2\pi + \frac{1}{2\pi} \frac{I_x}{\Delta} \frac{1}{\sigma^4} \stackrel{!}{=} 0 \\ \Leftrightarrow \sum_{j=1}^M \frac{4\pi}{\sigma^2} &= \sum_{j=1}^M \frac{1}{2\pi} \frac{I_x}{\Delta} \frac{1}{\sigma^4} \\ \Rightarrow \sigma^2 &= 2\pi \frac{1}{M} \sum_{j=1}^M I_x \frac{1}{\Delta} \end{aligned}$$

with  $\Delta = (1 + b^2 + 2b\cos(\lambda))$  and  $I_x$  as defined above. One can insert this into (\*) which is now only a function of  $b$  and will be maximised over  $b$  only. Inserting  $\sigma^2$  leads to cancellation of the  $2\pi$  in both cases so from now on  $\sigma^2 = \frac{1}{M} \sum_{j=1}^M I_x \frac{1}{\Delta}$ . Pulling apart the  $\ln$  then yields:

$$I_M(b) = \sum_{j=1}^M \left[ -\ln(\Delta) - \ln(\sigma^2) - \frac{I_x}{\Delta} \frac{1}{\sigma^2} \right] \quad (+)$$

## Coding Part

Here I define the functions that eventually make up the objective function (+) which I maximise. I tried to vectorize every function as far as possible to make the processing time somewhat shorter.

Getting the  $\Delta$  as describes above is achieved by running

```
delta <- function(b, lambdaj){  
  val <- (1 + b**2 + 2*b*cos(lambdaj))  
  return(val)  
}
```

This function can also take in the whole vector containing all  $\lambda_j$  and then returns an array of equal size with the corresponding  $\Delta$ s

The individual  $I_x$  are a function of  $\lambda_j$  and the provided time series  $x$ . This function takes the vector  $x$  and generated two vectors with the *sin* and *cos* respectively. Then I multiply elementwise and at the end take the sum over the whole vector.

```
Ix <- function(lambdaj, x){  
  
  bigT <- length(x)  
  little_t_seq <- seq(bigT)-1  
  
  prefactor <- 1/(2*pi*bigT)  
  
  sequence_sin <- cos(lambdaj * little_t_seq)  
  sequence_cos <- cos(lambdaj * little_t_seq)  
  
  val <- prefactor*(sum(x*sequence_cos)**2 + sum(x*sequence_sin)**2)  
  return(val)  
}
```

Coming to  $\sigma^2$ , the corresponding function takes as inputs the list of  $\lambda_j$ , a specific  $b$  and  $x$ . Then again *sumoverIx* and *sumoverdelta* are vectors of length  $M$  which I sum at the end.

```
sigma <- function(lambda_list, b, x){  
  
  bigM <- length(lambda_list)  
  
  sumoverIx <- sapply(lambda_list, Ix, x=x)  
  sumoverdelta <- delta(b=b, lambda_list)  
  
  val <- (1/bigM) * sum((sumoverIx/sumoverdelta))  
  return(val)  
}
```

The objective function (+) is finally given by

```
objf <- function(b, lambda_list, x){  
  
  firstterm <- delta(b, lambda_list)  
  secondterm <- sigma(lambda_list, b, x)  
  
  thirdterm_Ix <- sapply(lambda_list, Ix, x)  
  
  tmp <- -log(firstterm) - log(secondterm) - (thirdterm_Ix/firstterm)*(1/secondterm)
```

```

val <- sum(tmp)

return(val)
}

```

## Maximising $I_M$

The maximisation is then done by first defining a grid with  $\hat{b} \in (-1, 1)$  since if  $|\hat{b}| > 1$  the process would not be stable. For this I run the following loop where for first the data is taken from the dictionaries I computed before.

I find the optimal  $b_{opt}$  by maximising the objective function over the grid defined above. For calculating the corresponding  $\sigma^2$  I plug the  $b_{opt}$  into the *sigma* function and then multiply it by  $2\pi$  to get the correct value. The MSE of the estimated parameters  $b_{opt}$  and  $\sigma_{opt}$  is given by taking their respective biases to the power of two and adding the variance.

```

for (model in models){
  pos <- which(models==model)

  data_used <- data[[model]]
  sample_size <- dim(data_used)[1]
  ts_length <- dim(data_used)[2]

  df <- data.frame()

  for (s in seq(sample_size)){

    res <- sapply(b_list, objf, lambda_list=lambdas, x=data_used[s,])

    b_opt <- b_list[which.max(res)]
    sigma_opt <- sigma(lambdas, b_opt, x=data_used[s,]) * 2 * pi

    df <- rbind(df, cbind(b_opt, sigma_opt))
  }
  colnames(df) <- c("b", "sigma")

  bias_b <- (mean(df$b) - b)*ts_length
  bias_sigma <- (mean(df$sigma) - sigma_list[pos])*ts_length

  # Getting the values for MSE
  sigma_b <- var(df$b)
  sigma_sigma <- var(df$sigma)

  df_final <- rbind(df_final, cbind(bias_b, bias_sigma, sigma_b, sigma_sigma))
}

```

## Time Domain

### Theoretical Setup

First getting the conditional mean and variance of  $x_{t+1}$ :

$$\begin{aligned}
E[x_{t+1}|I_t] &= E[e_{t+1} + be_t|I_t] \\
&= E[e_{t+1}|I_t] + E[be_t|I_t] \\
&= be_t
\end{aligned}$$

and

$$\begin{aligned}
Var(x_{t+1}|I_t) &= E[x_{t+1}^2|I_t] - E[x_{t+1}|I_t]^2 \\
&= E[e_{t+1}^2 + 2be_{t+1}e_t + b^2e_t^2|I_t] - b^2e_t^2 \\
&= e_{t+1}^2 = \sigma^2
\end{aligned}$$

The conditional Loglikelihood is then given by

$$\begin{aligned}
LLH^c(b, \sigma^2) &= \sum_{t=1}^T \ln \left( \frac{1}{\sqrt{2\pi\sigma^2}} \right) - \frac{1}{2\sigma^2} (x_t - E[x_t|I_t])^2 \\
&= (T-1) \ln \left( \frac{1}{\sqrt{2\pi\sigma^2}} \right) - \frac{1}{2\sigma^2} \sum_{t=1}^T e_t^2
\end{aligned}$$

Since I only observe  $x_t$  I assume that  $e_0 = 0$  and approximate  $e_t$  by  $\hat{e}_t$

$$\begin{aligned}
\hat{e}_0 &= 0 \\
\hat{e}_1 &= x_1 \\
\hat{e}_2 &= x_2 - bx_1 \\
\hat{e}_3 &= x_3 - bx_2 \\
&\vdots \\
\hat{e}_t &= \sum_{s=0}^t x_{t-s} b^s (-1)^s
\end{aligned} \tag{-}$$

As in the frequency domain, I concentrate  $\sigma^2$  out such that the objective function to be maximised is only depending on one parameter,  $b$ .

$$\begin{aligned}
\frac{\partial LLH^c}{\partial \hat{\sigma}^2} &= \sum_{t=1}^T \sqrt{2\pi\hat{\sigma}^2} \left( -\frac{1}{2\sqrt{2\pi}} \frac{1}{\hat{\sigma}^2 \sqrt{\hat{\sigma}^2}} \right) + \left( \frac{1}{2\hat{\sigma}^4} \hat{e}_t^2 \right) \stackrel{!}{=} 0 \\
&\Leftrightarrow \sum_{t=1}^T \frac{1}{2\hat{\sigma}^2} = \sum_{t=1}^T \frac{1}{2\hat{\sigma}^4} \hat{e}_t^2 \\
&\Leftrightarrow \hat{\sigma}^2 = \frac{1}{T-1} \sum_{t=1}^T \hat{e}_t^2
\end{aligned}$$

For all three models (*i, ii, iii*) I assume the PDF of a normal distribution to see how the estimator compares to the one from the spectral density when having different distributions of the errors.

## Coding

For this part I switched to Python and vectorized everything which in the end was faster by a factor of 10 compared to R.

As a start I calculate  $(-)$ . I represent the  $x_{t-s}$  by a lower triangular matrix where the first column only contains  $x_0$  and then only zero. Column 2 then has  $\begin{pmatrix} x_1 \\ x_0 \end{pmatrix}$  followed by zeros and so on.

```
def make_xmatrix(x):  
  
    n = len(x)  
    xmat = np.zeros((n,n))  
  
    for i in range(n):  
  
        xtmp = x[0:n-i]  
  
        tmp = np.zeros(n)  
        tmp[n-len(xtmp):] = xtmp  
  
        xmat[i,:] = tmp  
    return xmat
```

The corresponding vector of  $b$  with the alternating signs is coded by

```
def make_bvector(b, x):  
    n = len(x)  
    bvector = np.repeat(b, n) ** np.arange(n)  
    bvector = bvector * ((-1) ** np.arange(n))  
  
    return bvector
```

I multiply that matrix and the vector elementwise to get a matrix of  $\hat{e}$ . Taking the sum over the individual columns then results in the vector of  $\hat{e}$ :

```
def make_ehat_mat(b, x):  
    n=len(x)  
  
    bvec = make_bvector(b, x)  
    xmat = make_xmatrix(x)  
  
    ehat_mat = np.empty((n,n))  
    for i in np.arange(n):  
        ehat_mat[:,i] = np.multiply(xmat[:,i], bvec)  
    return ehat_mat  
  
def make_ehat(ehatmat):  
  
    n = ehatmat.shape[1]  
    res = np.empty(n)  
    for i in range(n):  
        res[i] = np.sum(ehatmat[:,i])  
    return res
```

With this vector of  $\hat{e}_t$  the corresponding  $\hat{\sigma}^2$  are calculated with

```

def get_sigmahat(b,x):
    n = len(x)

    ehatmat = make_ehat_mat(b,x)
    eh = make_ehat(ehatmat)

    sigma = (1/(n-1)) * np.sum((eh**2))
    return sigma

```

The objective function is given by *objf\_MLE* and the optimisation is done by the function called *calc*.

```

def objf_MLE(b, x):
    n = len(x)
    sigmahat = get_sigmahat(b,x)
    logpart = (1/np.sqrt(2*np.pi*sigmahat))

    ehatmat = make_ehat_mat(b,x)
    eh = make_ehat(ehatmat)

    value = (n-1) * np.log(logpart) - (1/(2*sigmahat)) * np.sum((eh**2))
    return value

def calc(blist, x):
    samplesize = np.shape(x)[0]
    res_df = np.empty((samplesize, 2))

    for j in np.arange(samplesize):
        data_used = x[j,:]

        res_tmp = np.empty(len(blist))

        for i in np.arange(len(blist)):

            b_tmp = blist[i]
            res_tmp[i] = objf_MLE(b_tmp, data_used)

        b_opt = blist[np.where(np.max(res_tmp)==res_tmp)][0]
        sigma_opt = get_sigmahat(b_opt, data_used)

        res_df[j,:] = [b_opt, sigma_opt]
    return res_df

```

## Results

The values are then:

Table 1: Values from the frequency domain

	bias of $\hat{b}$	bias of $\hat{\sigma}^2$	MSE $\hat{b}$	MSE $\hat{\sigma}^2$
model <i>i</i>	4.40	-1.97	2.06	3.96
model <i>ii</i>	0.60	1.71	2.14	28.55
model <i>iii</i>	-1.55	-0.22	2.11	0.02

Table 2: Values from the time domain

	bias of $\hat{b}$	bias of $\hat{\sigma}^2$	MSE $\hat{b}$	MSE $\hat{\sigma}^2$
model <i>i</i>	3.10	-0.65	1.22	2.05
model <i>ii</i>	-0.05	3.1	1.17	22.21
model <i>iii</i>	404	167	163	27.95

In model *i* where the errors are drawn from a standard normal distribution the MLE estimator performs better than the frequency domain estimation. The same is still roughly true for the Student's *t*-distribution which is pretty similar to the standard normal.

As soon as the errors are drawn from a distribution substantially different as assumed in the MLE, this estimator performs very poorly in contrast to the frequency domain estimation. This estimator's performance does not depend on the distribution of the errors.

Suming up, when the distribution of the errors is known, the MLE estimator does a better job than using the Whittle estimator which still performs fairly good. With the distribution of the errors deviating from the assumed one, the MLE estimation should not be the preferred one. As most things in life, there is a trade-off striking again. Going for the more exact MLE estimation is coming with the drawbacks of having a really bad estimator when choosing the wrong underlying distribution. The Whittle estimator on the other hand is independent of the distribution but also less accurate compared to the optimal MLE.

## Question 16

### DGP

#### Preliminaries

For generating the data, I resort to using the following formula

$$y_t = \sum_{j=0}^{\infty} \psi_j \epsilon_{t-j} \quad (1)$$

with  $\psi_j = \frac{\Gamma(j+d)}{\Gamma(j+1)\Gamma(d)}$

and  $\Gamma(z) = \int_0^{\infty} \underbrace{x^{z-1} e^{-x}}_{f(x,z)} dx, \quad \text{where } \text{Re}(z) > 0$

However I approximate  $\psi_j$  by

$$\psi_j \sim \frac{j^{d-1}}{\Gamma(d)}$$

Since I cannot really take the sum in (1) to  $\infty$ , I take it up to  $M+T$  and will later discard the first  $M$  observations. here I code the function  $f$  and then the integral over that function. Taking  $d = 0$  makes it impossible to start at the lower bound of zero since then I'd be dividing by zero. This can be seen when taking the antiderivative of  $f$ . That's why I start with the lwoest possible value which is 1e-3 in that case.

```
f <- function(x, z){
  value <- x**(z-1)*exp(-x)
  return(value)
}
```



```

}

intf <- function(z){
  if (z != 0){
    value <- integrate(f, lower=0,upper=Inf, z=z)$value
  }else{
    value <- integrate(f, lower=1e-3,upper=Inf, z=z)$value
  }

  return(value)
}

```

For getting the  $\psi_j$  one needs to be careful with the order of the signs and the expression changes depending on  $d$  being negative or positive.

```

psi <- function(j, d){

  if (d<0){
    upper <- j**(-d-1)
    lower <- intf(abs(d))
  }else{
    upper <- j**(d-1)
    lower <- intf(d)
  }

  value <- upper/lower
  return(value)
}

```

## Getting the data

The final data gathering then happens according to the following function

```

get_data <- function(overall_length, d_list){
  y <- matrix(nrow=overall_length, ncol=length(d_list))

  for (i in seq_along(d_list)){
    # Initiating space
    res <- rep(NA, length(overall_length))

    # Summing
    for (t in seq_len(overall_length)){
      val <- 0
      for (j in seq_len(t-1)){
        tmp <- psi(j,d_list[i]) * et[t-j]
        val <- val + tmp
      }
      res[t] <- sum(val)
    }

    y[,i] <- res
  }
  return(y)
}

```

## Autocorrelations

For the autocorrelations I follow the slides and use

$$\gamma(0) = \sigma^2 \frac{\Gamma(1-2d)}{\Gamma^2(1-d)}$$
$$\gamma(h) = \frac{h-1+d}{h-d} \gamma(h-1)$$

Getting  $\gamma(0)$  is done by running

```
get_gamma0 <- function(data, d_list, d){  
  
  pos <- which(d_list==d)  
  sigma2 <- var(data[,pos])  
  upper <- intf(1-2*d)  
  lower <- intf(1-d)**2  
  
  final <- sigma2*(upper/lower)  
  return(final)  
}
```

which is a function of all of the  $d$  considered, the actual  $d$  as well as the data. Somewhat clunky implementation but works for now. The variance is just the estimated variance of the time series. \ The rest of the covariances and then subsequently the autocorrelations are given by this rather messy function

```
get_corrs <- function(data, laglength, d_list){  
  
  sto = dict()  
  for (i in seq_along(d_list)){  
    # Initiating storage space  
    lagframe <- data.frame()  
  
    # Getting gamma(0) and building up  
    gamma0 <- get_gamma0(data, d_list, d_list[i])  
    lagframe <- cbind(0, gamma0)  
  
    for (h in seq(laglength)){  
      pre <- (h-1+d_list[i])/(h-d_list[i])  
      gammah <- pre * lagframe[h,2]  
  
      lagframe <- rbind(lagframe, cbind(h, gammah))  
    }  
  
    # Calculating the autocorrelation from the covariances  
    lagframe[,2] <- lagframe[,2]/lagframe[1,2]  
    colnames(lagframe) <- c("h", 'real')  
  
    # Getting the estimated autocorrelations  
    estimated <- acf(data[,i], lag.max=laglength,  
                     plot=FALSE, demean = FALSE)$acf  
    final <- cbind(lagframe, estimated)
```


```

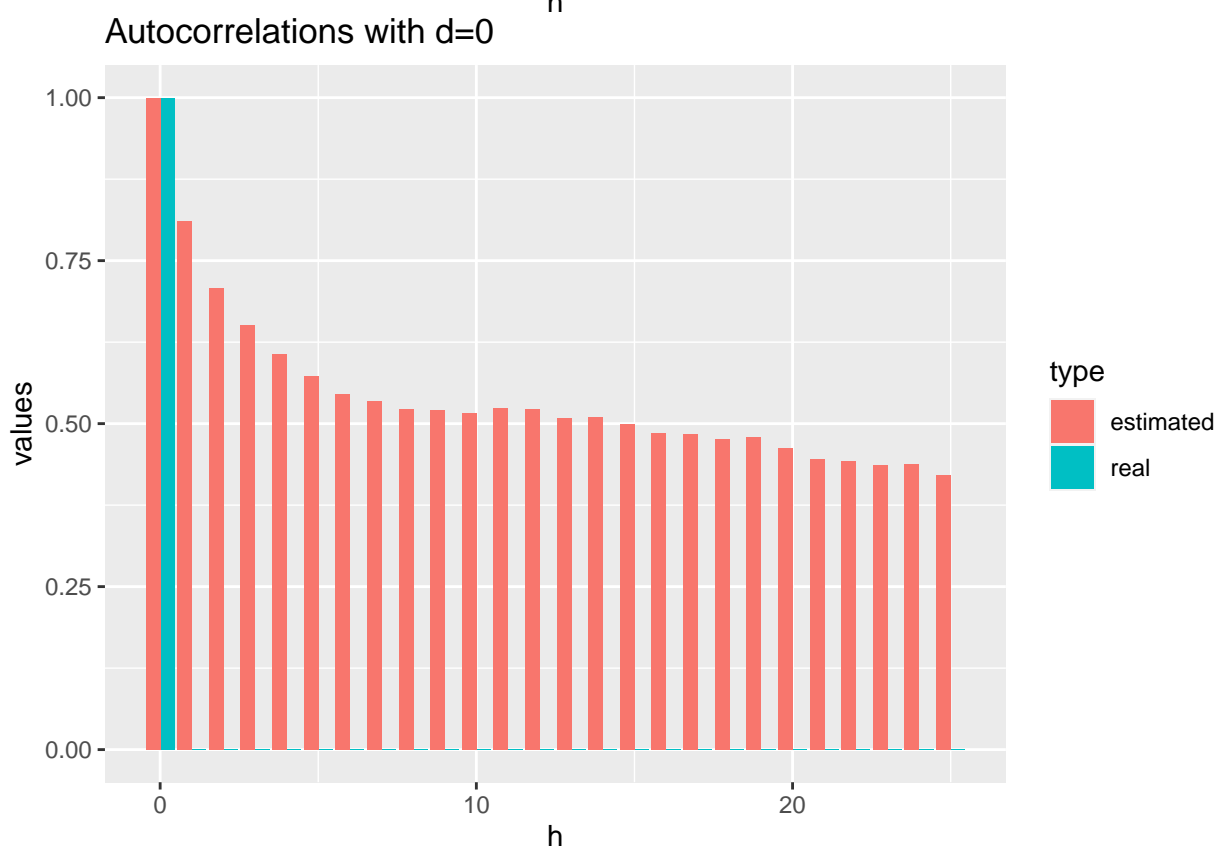
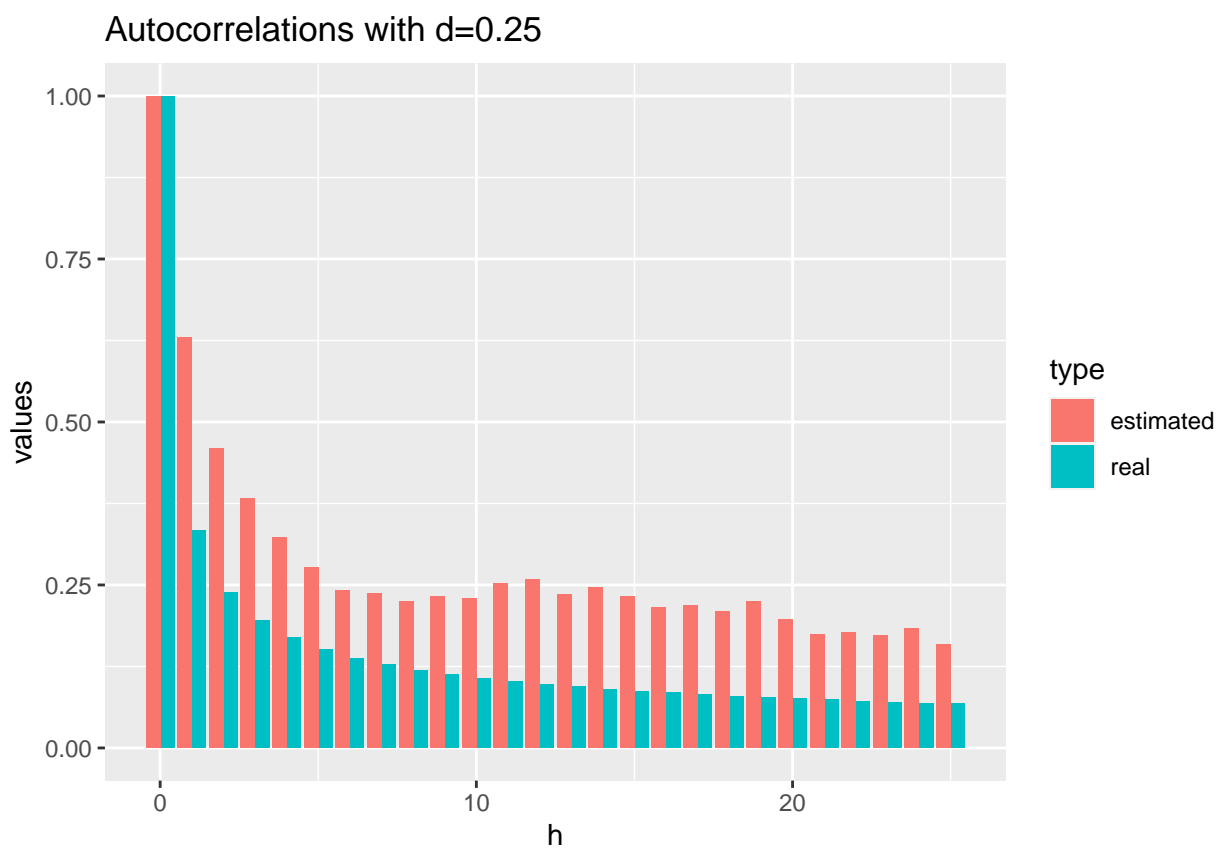
    sto[[toString(i)]] <- data.frame(final)
  }


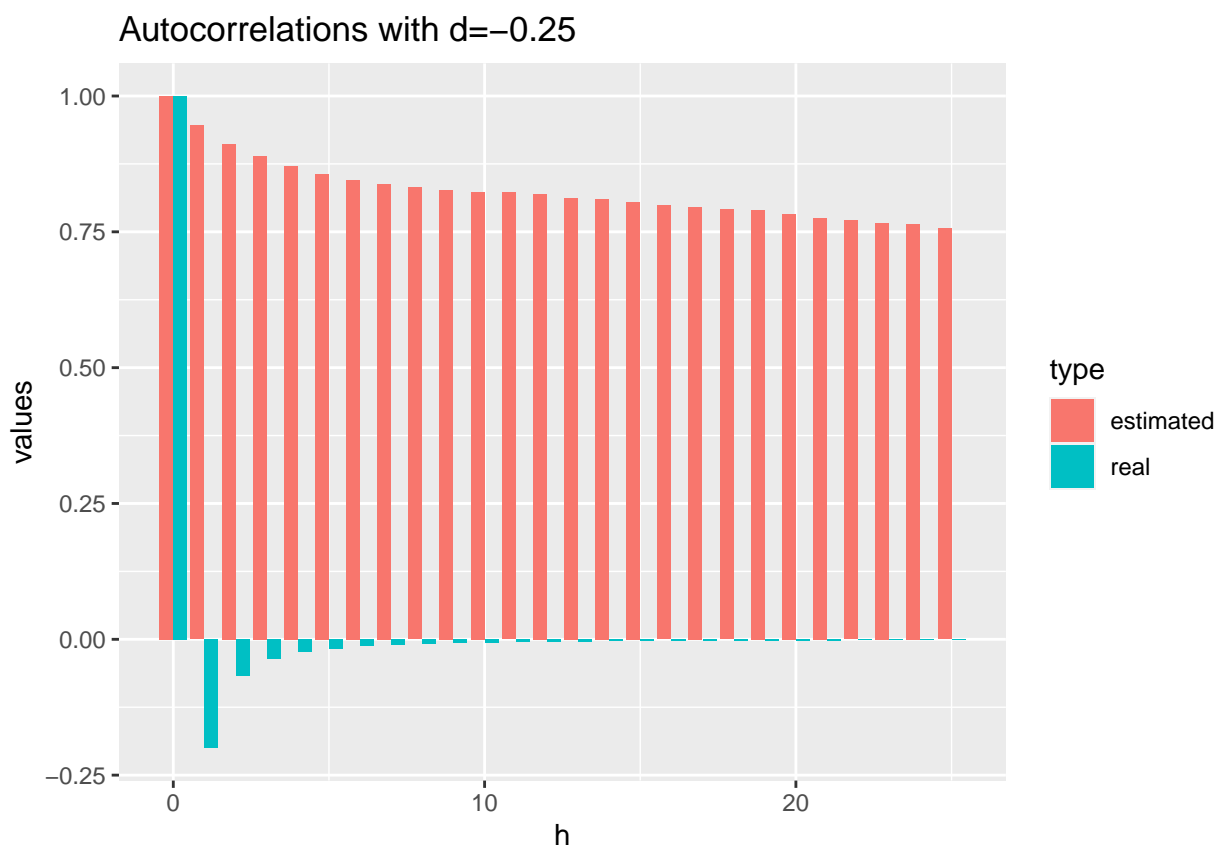
  return(sto)
}

```

It also returns the estimated autocorrelations up to lag 25 which for which I used the standard R function *acf*. This is then saved in a dictionary. The results depending on the different  $d$  are given in the figures below. The values for *real* are calculated with the functions described above while the values for *estimated* come from the *acf* function which estimates the autocorrelations of a given time series. The very first value is 1 all the time since it is the zero-th lag which simply  $\frac{\gamma(0)}{\gamma(0)} = 1$ .









Now the plots with  $y_t$ . The interesting thing here is that the time series for  $d = -0.45$  and  $d = 0.45$  is basically

the same, in the second case you almost cannot see it since it's overlayed by the other one. Apparently it does not make a difference to the time series whether it is fractally integrated by amount  $\delta$  or  $-\delta$ . This can also be seen when deriving and coding the DGP. I always switch the sign such that the expression evaluated is the same.


Time series of y with d=0.45




Time series of y with d=0.25




Time series of y with d=0



Time series of y with  $d=-0.25$



Time series of y with  $d=-0.45$





## Appendix

All other functions used for both questions:

```
prepplot <- function(data){  
  df_plot <- pivot_longer(data, -h, names_to = 'type', values_to = 'values')  
  return(df_plot)  
}  
  
shift <- function(x, lag) {  
  n <- length(x)  
  xnew <- rep(NA, n)  
  if (lag < 0) {  
    xnew[1:(n-abs(lag))] <- x[(abs(lag)+1):n]  
  } else if (lag > 0) {  
    xnew[(lag+1):n] <- x[1:(n-lag)]  
  } else {  
    xnew <- x  
  }  
  return(xnew)  
}
```

```
def get_biases(dataf, model):  
  bias_b = (np.mean(dataf[:,0]) - b)*1000  
  bias_sigma = (np.mean(dataf[:,1]) - sigma_list[model-1])*1000  
  
  sigma_b = np.var(dataf[:,0])  
  sigma_sigma = np.var(dataf[:,1])  
  
  MSE_b = ((bias_b/1000)**2 + sigma_b) * 1000  
  MSE_sigma = ((bias_sigma/1000)**2 + sigma_sigma) * 1000  
  
  return [bias_b , bias_sigma, MSE_b, MSE_sigma]
```